

Start coding or [generate](#) with AI.

Write a python program to compute Mean, Median, Mode, Variance, Standard Deviation using the Iris dataset. Also, demonstrate various data pre-processing techniques for a random dataset, including reshaping, filtering, merging, handling missing values, and Min-max normalization.

✓ Load the iris dataset

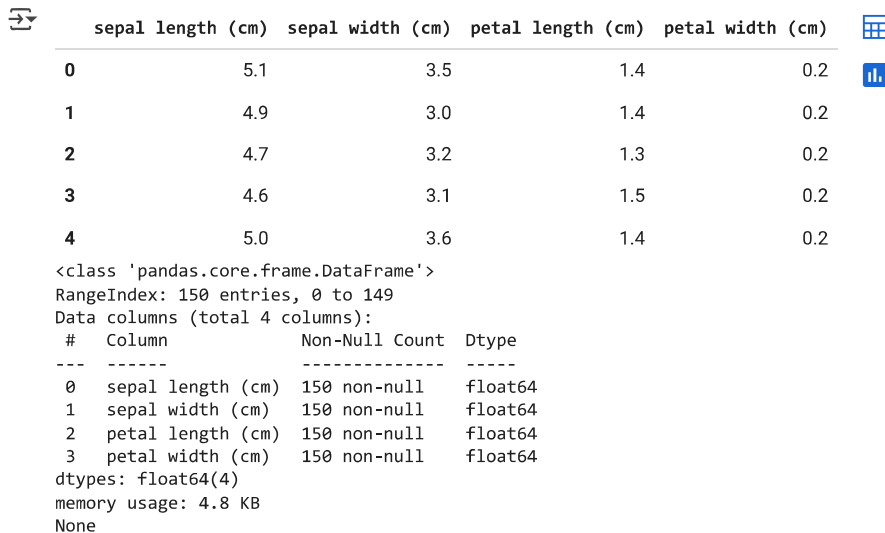
Load the Iris dataset using a suitable library like scikit-learn or pandas.

```
from sklearn.datasets import load_iris
import pandas as pd

iris = load_iris()

df_iris = pd.DataFrame(data=iris.data, columns=iris.feature_names)

display(df_iris.head())
display(df_iris.info())
```



	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sepal length (cm)      150 non-null   float64
1   sepal width (cm)       150 non-null   float64
2   petal length (cm)      150 non-null   float64
3   petal width (cm)       150 non-null   float64
dtypes: float64(4)
memory usage: 4.8 KB
None
```

✓ Compute descriptive statistics

Calculate the Mean, Median, Mode, Variance, and Standard Deviation for each feature in the Iris dataset.

```
import numpy as np

mean_values = df_iris.mean()

median_values = df_iris.median()

mode_values = df_iris.mode().iloc[0]


variance_values = df_iris.var()

std_dev_values = df_iris.std()




statistics = {
    'Mean': mean_values,
    'Median': median_values,
    'Mode': mode_values,
    'Variance': variance_values,
    'Standard Deviation': std_dev_values
}

stats_df = pd.DataFrame(statistics)

display(stats_df)
```



	Mean	Median	Mode	Variance	Standard Deviation
sepal length (cm)	5.843333	5.80	5.0	0.685694	0.828066
sepal width (cm)	3.057333	3.00	3.0	0.189979	0.435866
petal length (cm)	3.758000	4.35	1.4	3.116278	1.765298
petal width (cm)	1.199333	1.30	0.2	0.581006	0.762238




Next steps: [Generate code with stats_df](#) [View recommended plots](#) [New interactive sheet](#)




Display results

Present the computed statistics in a clear format.

```
display(stats_df)
```



	Mean	Median	Mode	Variance	Standard Deviation
sepal length (cm)	5.843333	5.80	5.0	0.685694	0.828066
sepal width (cm)	3.057333	3.00	3.0	0.189979	0.435866
petal length (cm)	3.758000	4.35	1.4	3.116278	1.765298
petal width (cm)	1.199333	1.30	0.2	0.581006	0.762238



Next steps: [Generate code with stats_df](#) [View recommended plots](#) [New interactive sheet](#)

Create a random dataset

Generate a random dataset for demonstration purposes.

```
import numpy as np
import pandas as pd

data = {
    'Column A': np.random.rand(15),
    'Column B': np.random.rand(15),
    'Column C': np.random.rand(15)
}

df_random = pd.DataFrame(data)

display(df_random.head())
```



	Column A	Column B	Column C
0	0.465538	0.445022	0.759710
1	0.135850	0.271254	0.252382
2	0.467719	0.046016	0.628122
3	0.174865	0.394590	0.240067
4	0.581747	0.375265	0.485304




Reshape the data


Demonstrate reshaping the data using appropriate techniques (e.g., melt, pivot).

```
df_melted = pd.melt(df_random, var_name='variable', value_name='value')

display(df_melted.head())
```




	variable	value
0	Column A	0.465538
1	Column A	0.135850
2	Column A	0.467719
3	Column A	0.174865
4	Column A	0.581747




Filter the data

Show how to filter the data based on specific conditions.

```
df_filtered = df_random[df_random['Column A'] > 0.5]
display(df_filtered.head())
```



	Column A	Column B	Column C
4	0.581747	0.375265	0.485304
9	0.731315	0.702102	0.119362
10	0.993226	0.782158	0.621174
11	0.959279	0.774722	0.030088
12	0.556527	0.852705	0.899067



Merge the data


Create another random dataset and demonstrate merging it with the first dataset.

```
data2 = {
    'Column D': np.random.rand(len(df_random))
}


df_random2 = pd.DataFrame(data2)

df_merged = pd.concat([df_random, df_random2], axis=1)

display(df_merged.head())
```



	Column A	Column B	Column C	Column D
0	0.465538	0.445022	0.759710	0.010975
1	0.135850	0.271254	0.252382	0.343383
2	0.467719	0.046016	0.628122	0.632905
3	0.174865	0.394590	0.240067	0.125320
4	0.581747	0.375265	0.485304	0.779490




Handle missing values

Introduce missing values into the dataset and demonstrate techniques for handling them (e.g., imputation, dropping).

```
import numpy as np

np.random.seed(42)
mask = np.random.choice([True, False], size=df_merged.shape, p=[0.1, 0.9])
df_merged_with_missing = df_merged.mask(mask)

display("DataFrame with missing values:")
display(df_merged_with_missing)
```

 'DataFrame with missing values:'

	Column A	Column B	Column C	Column D
0	0.465538	0.445022	0.759710	0.010975
1	0.135850	0.271254	NaN	0.343383
2	0.467719	0.046016	NaN	0.632905
3	0.174865	0.394590	0.240067	0.125320
4	0.581747	0.375265	0.485304	0.779490
5	0.117306	0.770830	0.030528	0.020289
6	0.303839	0.100768	0.303682	0.067747
7	0.324920	NaN	0.489434	0.435235
8	NaN	0.980674	0.688351	0.462965
9	0.731315	NaN	0.119362	0.583026
10	0.993226	0.782158	NaN	0.433757
11	0.959279	0.774722	0.030088	0.219863
12	0.556527	0.852705	0.899067	0.047146
13	0.073683	0.234956	0.381687	0.999222
14	NaN	0.537056	NaN	0.049887

Next steps:

[Generate code with df_merged_with_missing](#)

[View recommended plots](#)


[New interactive sheet](#)

```
df_dropped = df_merged_with_missing.dropna()

display("DataFrame after dropping rows with missing values (head):")
display(df_dropped.head())

df_imputed = df_merged_with_missing.fillna(df_merged_with_missing.mean())

display("DataFrame after imputing missing values with mean (head):")
display(df_imputed.head())
```

 'DataFrame after dropping rows with missing values (head):'

	Column A	Column B	Column C	Column D
0	0.465538	0.445022	0.759710	0.010975
3	0.174865	0.394590	0.240067	0.125320
4	0.581747	0.375265	0.485304	0.779490
5	0.117306	0.770830	0.030528	0.020289
6	0.303839	0.100768	0.303682	0.067747

'DataFrame after imputing missing values with mean (head):'

	Column A	Column B	Column C	Column D
0	0.465538	0.445022	0.759710	0.010975
1	0.135850	0.271254	0.402480	0.343383
2	0.467719	0.046016	0.402480	0.632905
3	0.174865	0.394590	0.240067	0.125320
4	0.581747	0.375265	0.485304	0.779490

Perform min-max normalization

Apply Min-max normalization to a numerical feature in the dataset.

```
column_to_normalize = 'Column A'

min_val = df_imputed[column_to_normalize].min()
```


```

max_val = df_imputed[column_to_normalize].max()


df_imputed[f'{column_to_normalize}_normalized'] = (df_imputed[column_to_normalize] - min_val) / (max_val - min_val)

display(df_imputed.head())

```



	Column A	Column B	Column C	Column D	Column A_normalized
0	0.465538	0.445022	0.759710	0.010975	0.426141
1	0.135850	0.271254	0.402480	0.343383	0.067606
2	0.467719	0.046016	0.402480	0.632905	0.428513
3	0.174865	0.394590	0.240067	0.125320	0.110034
4	0.581747	0.375265	0.485304	0.779490	0.552517



✓ Display results

Show the results of each pre-processing step.

```

display("Original Random DataFrame:")
display(df_random)

display("Melted DataFrame:")
display(df_melted)

display("Filtered DataFrame:")
display(df_filtered)


display("Merged DataFrame:")
display(df_merged)

display("DataFrame with Missing Values:")
display(df_merged_with_missing)




display("DataFrame after Dropping Missing Values:")
display(df_dropped)

display("DataFrame after Imputing Missing Values and Normalization:")
display(df_imputed)

```



 'Original Random DataFrame: '

	Column A	Column B	Column C
0	0.465538	0.445022	0.759710
1	0.135850	0.271254	0.252382
2	0.467719	0.046016	0.628122
3	0.174865	0.394590	0.240067
4	0.581747	0.375265	0.485304
5	0.117306	0.770830	0.030528
6	0.303839	0.100768	0.303682
7	0.324920	0.558833	0.489434
8	0.497950	0.980674	0.688351
9	0.731315	0.702102	0.119362
10	0.993226	0.782158	0.621174
11	0.959279	0.774722	0.030088
12	0.556527	0.852705	0.899067
13	0.073683	0.234956	0.381687
14	0.661941	0.537056	0.756287





'Melted DataFrame: '

	variable	value
0	Column A	0.465538
1	Column A	0.135850
2	Column A	0.467719
3	Column A	0.174865
4	Column A	0.581747
5	Column A	0.117306
6	Column A	0.303839
7	Column A	0.324920
8	Column A	0.497950
9	Column A	0.731315
10	Column A	0.993226
11	Column A	0.959279
12	Column A	0.556527
13	Column A	0.073683
14	Column A	0.661941
15	Column B	0.445022
16	Column B	0.271254
17	Column B	0.046016
18	Column B	0.394590
19	Column B	0.375265
20	Column B	0.770830
21	Column B	0.100768
22	Column B	0.558833
23	Column B	0.980674
24	Column B	0.702102
25	Column B	0.782158
26	Column B	0.774722
27	Column B	0.852705
28	Column B	0.234956





```
28 Column B 0.234956
29 Column B 0.537056
30 Column C 0.759710
31 Column C 0.252382
32 Column C 0.628122
33 Column C 0.240067
34 Column C 0.485304
35 Column C 0.030528
36 Column C 0.303682
37 Column C 0.489434
38 Column C 0.688351
39 Column C 0.119362
40 Column C 0.621174
41 Column C 0.030088
42 Column C 0.899067
43 Column C 0.381687
44 Column C 0.756287
```



'Filtered DataFrame:'

	Column A	Column B	Column C	
4	0.581747	0.375265	0.485304	
9	0.731315	0.702102	0.119362	
10	0.993226	0.782158	0.621174	
11	0.959279	0.774722	0.030088	
12	0.556527	0.852705	0.899067	
14	0.661941	0.537056	0.756287	

'Merged DataFrame:'

	Column A	Column B	Column C	Column D	
0	0.465538	0.445022	0.759710	0.010975	
1	0.135850	0.271254	0.252382	0.343383	
2	0.467719	0.046016	0.628122	0.632905	
3	0.174865	0.394590	0.240067	0.125320	
4	0.581747	0.375265	0.485304	0.779490	
5	0.117306	0.770830	0.030528	0.020289	
6	0.303839	0.100768	0.303682	0.067747	
7	0.324920	0.558833	0.489434	0.435235	
8	0.497950	0.980674	0.688351	0.462965	
9	0.731315	0.702102	0.119362	0.583026	
10	0.993226	0.782158	0.621174	0.433757	
11	0.959279	0.774722	0.030088	0.219863	
12	0.556527	0.852705	0.899067	0.047146	
13	0.073683	0.234956	0.381687	0.999222	
14	0.661941	0.537056	0.756287	0.049887	

'DataFrame with Missing Values:'

	Column A	Column B	Column C	Column D	
0	0.465538	0.445022	0.759710	0.010975	
1	0.135850	0.271254	NaN	0.343383	
2	0.467719	0.046016	NaN	0.632905	
3	0.174865	0.394590	0.240067	0.125320	
4	0.581747	0.375265	0.485304	0.779490	